

GARUDA 2.0

Sage Sefton	Avinash Karanth
ss557415@ohio.edu	karanth@ohio.edu
Ohio University	Ohio University

Gordon Stewart
gstewart@ohio.edu
Ohio University

2020
Spring

1 Background

GARUDA was a system that operated on a two stream model of the "user" or program, and that which executed the instructions. GARUDA 2.0, we extend, or refine based on the point of view, this model to operate between the Execution (EXE) and Memory (MEM) stages of the pipeline. We assume a monitor theory shown by the graphic Fig ?? . The monitor appears to operate on a single stream, but the two hanging connections can be considered the two streams of the former model.

The most obvious change we need to make is to enable encryption and decryption of the EffAddr. Our attack vector focuses on a trojan targeting the state register between the EXE and MEM stages. They could cache these addresses and try to access that memory later, or they may try to use them as a side channel. Either way, allowing an adversary to track our memory accesses is clearly problematic. However, when we refer to "encryption", we really mean some reversible obfuscation function. We employ the proof assistant, Coq, to prove the EN and DE blocks are exact inverses. GARUDA is written in Coq for exactly this reason, so requiring a proof that $\forall P, DE(ENP) = P$ should be sufficient.

2 Syntax

Instruction Fields	f_i	$::=$	$f_{i_1} \mid \cdots \mid f_{i_k}$
Result Fields	f_r	$::=$	$f_{r_1} \mid \cdots \mid f_{r_k}$
Fields	f	$::=$	$f_i \mid f_r$
Instructions	i	$::=$	$\{f_{i_1} = v_{i_1}, \dots, f_{i_k} = v_{i_k}\}$
Results	r	$::=$	$\{f_{r_1} = v_{r_1}, \dots, f_{r_k} = v_{r_k}\}$
Predicates	a, b	$::=$	0 <i>False</i>
			1 <i>Identity</i>
			$f = n$ <i>Test</i>
			$a + b$ <i>Sum</i>
			$a \cdot b$ <i>Product</i>
			$\neg a$ <i>Negation</i>

Policies	p, q	$::=$	test (a)	<i>Test</i>
			$act(p)$	<i>Slice Actions</i>
			$res(p)$	<i>Slice Results</i>
			inj_i	<i>Injection Action</i>
			inj_r	<i>Injection Result</i>
			$f \leftarrow n$	<i>Update</i>
			$p + q$	<i>Choice</i>
			$p \cdot q$	<i>Sequential Concatenation</i>

3 Semantics

$$\begin{aligned}
\llbracket \cdot \rrbracket &: \mathbf{Stream}(I) \times \mathbf{Stream}(R) \rightarrow \\
&\quad P(\mathbf{Stream}(I)) \times P(\mathbf{Stream}(R)) \\
\llbracket 0 \rrbracket(-, -) &\triangleq (\emptyset, \emptyset) \\
\llbracket 1 \rrbracket(is, rs) &\triangleq (\{is\}, \{rs\}) \\
\llbracket f = n \rrbracket(is, rs) &\triangleq (\text{filter } (f = n) \{is\}, \text{filter } (f = n) \{rs\}) \\
\llbracket a + b \rrbracket(is, rs) &\triangleq \llbracket a \rrbracket(is, rs) \cup \llbracket b \rrbracket(is, rs) \\
&\quad \text{where } (S_i^1, S_r^1) \cup (S_i^2, S_r^2) \triangleq (S_i^1 \cup S_i^2, S_r^1 \cup S_r^2) \\
\llbracket a \cdot b \rrbracket(is, rs) &\triangleq \llbracket a \rrbracket(is, rs) \cap \llbracket b \rrbracket(is, rs) \\
&\quad \text{where } (S_i^1, S_r^1) \cap (S_i^2, S_r^2) \triangleq (S_i^1 \cap S_i^2, S_r^1 \cap S_r^2) \\
\llbracket \neg a \rrbracket(is, rs) &\triangleq \text{let } (S_i, S_r) = \llbracket a \rrbracket(is, rs) \\
&\quad \text{in } (\{is\} - S_i, \{rs\} - S_r) \\
\llbracket \text{act}(p) \rrbracket(is, rs) &\triangleq \text{let } (S_i, S_r) = \llbracket p \rrbracket(is, rs) \text{ in } (S_i, \{rs\}) \\
\llbracket \text{res}(p) \rrbracket(is, rs) &\triangleq \text{let } (S_i, S_r) = \llbracket p \rrbracket(is, rs) \text{ in } (\{is\}, S_r) \\
\llbracket \text{inj}_i(i) \rrbracket(is, rs) &\triangleq (\{i : is\}, \{rs\}) \\
\llbracket \text{inj}_r(r) \rrbracket(is, rs) &\triangleq (\{is\}, \{r : rs\}) \\
\llbracket f \leftarrow n \rrbracket(is, rs) &\triangleq (\text{map } (f \leftarrow n) \{is\}, \text{map } (f \leftarrow n) \{rs\}) \\
\llbracket p + q \rrbracket(is, rs) &\triangleq \llbracket p \rrbracket(is, rs) \cup \llbracket q \rrbracket(is, rs) \\
&\quad \text{where } (S_i^1, S_r^1) \cup (S_i^2, S_r^2) \triangleq (S_i^1 \cup S_i^2, S_r^1 \cup S_r^2) \\
\llbracket p \cdot q \rrbracket(is, rs) &\triangleq \text{let } (S_i, S_r) = \llbracket p \rrbracket(is, rs) \\
&\quad \text{in } \bigcup \{ \llbracket q \rrbracket(is', rs') \mid is' \in S_i, rs' \in S_r \} \\
\llbracket \text{filter } f \rrbracket(S) &\triangleq \{l \in S \mid f(l) = \text{true}\} \\
\llbracket \text{map } g \rrbracket(S) &\triangleq \{g(l) \mid l \in S\}
\end{aligned}$$

4 Intermediate Syntax

Values $v ::= \text{INSTR} \mid \text{RES}$

Registers $b ::= \text{reg}$

Expressions $e ::=$	$\text{read}(b)$	<i>Read Reg</i>
	$\mid \text{write}(b, v)$	<i>Write Value to Reg</i>
	$\mid \text{let } x = e_1 \text{ in } e_2$	<i>Assignment</i>
	$\mid e_1 \parallel e_2$	<i>Parallel</i>
	$\mid f(e_1)$	<i>Apply Function</i>
	$\mid \text{if } v = n \text{ then } e_1 \text{ else } e_2$	<i>Conditional</i>
	$\mid e_1 \ \&\& \ e_2$	<i>Product (AND)</i>
	$\mid e_1; e_2$	<i>Concatination</i>
	$\mid \text{forever } e$	<i>Hardwire</i>

5 Intermediate Semantics

Assume all \mathbf{C} is short for $\mathbf{C}_{(i_{in}, i_{out}, r_{in}, r_{out})}$ if unspecified.

$$\begin{aligned}
 \mathbf{C}[0] &= \text{let } i_bog = \text{new buf} \\
 &\quad r_bog = \text{new buf} \\
 &\quad \text{in } \text{write}(i_{out}, i_bog) \\
 &\quad \text{write}(r_{out}, r_bog) \\
 \mathbf{C}[1] &= \text{write}(i_{out}, \text{read}(i_{in})) \\
 &\quad \text{write}(r_{out}, \text{read}(r_{in})) \\
 \mathbf{C}[f = n] &= \text{if } i_{in} = n \text{ then} \\
 &\quad \text{write}(i_{out}, \text{read}(i_{in})) \\
 &\quad \text{else } \mathbf{C}_{(i_{in}, i_{out}, -, -)}[0] \\
 &\quad \text{if } r_{in} = n \text{ then} \\
 &\quad \text{write}(r_{out}, \text{read}(r_{in})) \\
 &\quad \text{else } \mathbf{C}_{(-, -, r_{in}, r_{out})}[0]
 \end{aligned}$$

$$\begin{aligned}
\mathbf{C}[\text{test}(a + b)] &= \text{let } e_{a_{ii}}, e_{a_{ri}}, e_{b_{ii}}, e_{b_{ri}} = \text{new buf in} \\
&\quad DeMux(i_{in}, e_{a_{ii}}, e_{b_{ii}}) \parallel DeMux(r_{in}, e_{a_{ri}}, e_{b_{ri}}) \\
&\quad \text{let } e_a = \mathbf{C}_{(e_{a_{ii}}, e_{a_{io}}, e_{a_{ri}}, e_{a_{ro}})}[a] \\
&\quad \quad e_b = \mathbf{C}_{(e_{b_{ii}}, e_{b_{io}}, e_{b_{ri}}, e_{b_{ro}})}[b] \\
&\quad \text{in } Mux(e_a \ e_b, (i_{out}, r_{out})) \\
\mathbf{C}[\text{test}(a \cdot b)] &= \mathbf{C}[\text{test}(a) \cdot \text{test}(b)] \\
\mathbf{C}[\neg a] &= \text{if } \neg a \text{ then} \\
&\quad \text{write}(i_{out}, \text{read}(i_{in})) \\
&\quad \text{write}(r_{out}, \text{read}(r_{in})) \\
&\quad \text{else } \mathbf{C}[0] \\
\mathbf{C}[\text{act}(p)] &= \text{let } r_bog_i = \text{new buf} \\
&\quad \quad r_bog_o = \text{new buf} \\
&\quad \text{in } e_p = \mathbf{C}_{(i_{in}, i_{out}, r_bog_i, r_bog_o)}[p] \\
&\quad \quad e_p \parallel \text{write}(r_{out}, \text{read}(r_{in})) \\
\mathbf{C}[\text{res}(p)] &= \text{let } i_bog_i = \text{new buf} \\
&\quad \quad i_bog_o = \text{new buf} \\
&\quad \text{in } e_p = \mathbf{C}_{(i_bog_i, i_bog_o, r_{in}, r_{out})}[p] \\
&\quad \quad e_p \parallel \text{write}(i_{out}, \text{read}(i_{in})) \\
\mathbf{C}[\text{inj}_i V] &= \text{write}(i_{out}, V) \\
\mathbf{C}[\text{inj}_r V] &= \text{write}(r_{out}, V) \\
\mathbf{C}[f \leftarrow n] &= \text{let } i' = f(\text{read}(i_{in})) \\
&\quad \quad r' = f(\text{read}(r_{in})) \\
&\quad \text{in } \text{write}(i_{out}, i') \\
&\quad \quad \text{write}(r_{out}, r') \\
\mathbf{C}[p + q] &= \text{let } e_p = \mathbf{C}[p] \\
&\quad \quad e_q = \mathbf{C}[q] \\
&\quad \text{in } Mux(e_p, e_q, (i_{out}, r_{out})) \\
\mathbf{C}[p \cdot q] &= \text{let } i_{mid} = \text{new buf} \\
&\quad \quad r_{mid} = \text{new buf} \\
&\quad \quad e_p = \mathbf{C}_{(i_{in}, i_{mid}, r_{in}, r_{mid})}[p] \\
&\quad \quad e_q = \mathbf{C}_{(i_{mid}, i_{out}, r_{mid}, r_{out})}[q] \\
&\quad \text{in } e_p ; e_q
\end{aligned}$$

6 Applications of GARUDA 2.0

6.1 Standard Taint

In the previous version of GARUDA, taint was implemented by tagging the most significant bit of a register. This caused no hardware overhead to maintain in the pipeline, but imposed an obvious bit-resolution hindrance. In GARUDA 2.0, the definition of such a policy is identical, but the compilation is different.

We assume a MIPS-like architecture for this example. The op codes of any given instruction are no more than 3; IN_1 , IN_2 , and OUT . We denote these as RS , RT , and RD .

Let's suppose you want your taint to propagate on the inputs of any arithmetic instructions. Additionally, you prohibit any tainted instructions or addresses from accessing memory. In GARUDA 2.0, we can define this as follows. Please note that ALU can further be expanded to specific types of ALU instructions

$$\begin{aligned}
 \text{Instruction Fields } f_i &::= \text{Taint}_{RS}, \text{Taint}_{RT}, \text{Taint}_{RD}, \text{OP} \\
 \text{OP} &::= \text{MEM}_{\text{READ}} \mid \text{MEM}_{\text{WRITE}} \mid \text{ALU} \mid \dots \\
 \text{Result Fields } f_r &::= (*\text{empty}*) \\
 \\
 \text{Arith} &\triangleq \text{OP} = \text{ALU} \\
 \text{Read} &\triangleq \text{OP} = \text{MEM}_{\text{READ}} \\
 \text{Write} &\triangleq \text{OP} = \text{MEM}_{\text{WRITE}} \\
 \text{AnyMem} &\triangleq \text{Read} + \text{Write} \\
 \\
 \text{TaintedInstr} &\triangleq (\text{Taint}_{RS} = \text{TRUE}) + (\text{Taint}_{RT} = \text{TRUE}) \\
 \text{TaintRes} &\triangleq \text{Taint}_{RD} \leftarrow \text{TRUE} \\
 \\
 \text{PropTaintALU} &\triangleq \text{act}(\text{Arith} \cdot \text{TaintedInstr} \cdot \text{TaintRes}) \\
 &\quad + \text{act}(\text{Arith} \cdot \neg \text{TaintedInstr}) \\
 \text{NoTaintMem} &\triangleq \text{act}(\text{AnyMem} \cdot \neg \text{TaintedInstr}) \\
 \text{SecureMem} &\triangleq \text{PropTaintALU} + \text{NoTaintMem} \\
 &\quad + \text{act}(\neg(\text{Arith} + \text{AnyMem}))
 \end{aligned}$$

6.2 Speculative Taint

[1]

$$\begin{aligned}
 \text{Instruction Fields } f_i &::= \text{Taint}_{\text{RS}}, \text{Taint}_{\text{RT}}, \text{Taint}_{\text{RD}}, \text{OP} \\
 \text{OP} &::= \text{MEM}_{\text{READ}} \mid \text{MEM}_{\text{WRITE}} \mid \text{ALU} \mid \dots \\
 \text{Result Fields } f_r &::= (*\text{empty}*) \\
 \\
 \text{Arith} &\triangleq \text{OP} = \text{ALU} \\
 \text{Read} &\triangleq \text{OP} = \text{MEM}_{\text{READ}} \\
 \text{Write} &\triangleq \text{OP} = \text{MEM}_{\text{WRITE}} \\
 \text{AnyMem} &\triangleq \text{Read} + \text{Write} \\
 \\
 \text{TaintedInstr} &\triangleq (\text{Taint}_{\text{RS}} = \text{TRUE}) + (\text{Taint}_{\text{RT}} = \text{TRUE}) \\
 \text{TaintRes} &\triangleq \text{Taint}_{\text{RD}} \leftarrow \text{TRUE} \\
 \\
 \text{PropTaintALU} &\triangleq \text{act}(\text{Arith} \cdot \text{TaintedInstr} \cdot \text{TaintRes}) \\
 &\quad + \text{act}(\text{Arith} \cdot \neg \text{TaintedInstr}) \\
 \text{NoTaintMem} &\triangleq \text{act}(\text{AnyMem} \cdot \neg \text{TaintedInstr}) \\
 \text{SecureMem} &\triangleq \text{PropTaintALU} + \text{NoTaintMem} \\
 &\quad + \text{act}(\neg(\text{Arith} + \text{AnyMem}))
 \end{aligned}$$

References

- [1] Jiyong Yu, Mengjia Yan, Artem Khyzha, Adam Morrison, Josep Torrellas, and Christopher W Fletcher. Speculative taint tracking (stt) a comprehensive protection for speculatively accessed data. In *Proceedings of the 52nd Annual IEEE/ACM International Symposium on Microarchitecture*, pages 954–968, 2019.